

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
df = pd.read_csv('processed.cleveland.data',names= ['age','sex','chest_pain','blood pressure',
'electrocardiographic','max_heart_rate','induced_angina','ST_depression','s
```

```
df
```

	age	sex	chest_pain	blood pressure	serum_cholesterol	fasting_blood_sugar	electro
<b>0</b>	63.0	1.0	1.0	145.0	233.0	1.0	
<b>1</b>	67.0	1.0	4.0	160.0	286.0	0.0	
<b>2</b>	67.0	1.0	4.0	120.0	229.0	0.0	
<b>3</b>	37.0	1.0	3.0	130.0	250.0	0.0	
<b>4</b>	41.0	0.0	2.0	130.0	204.0	0.0	
...	...	...	...	...	...	...	...
<b>298</b>	45.0	1.0	1.0	110.0	264.0	0.0	
<b>299</b>	68.0	1.0	4.0	144.0	193.0	1.0	
<b>300</b>	57.0	1.0	4.0	130.0	131.0	0.0	
<b>301</b>	57.0	0.0	2.0	130.0	236.0	0.0	
<b>302</b>	38.0	1.0	3.0	138.0	175.0	0.0	

303 rows × 14 columns

```
df.describe()
```

```

    age          sex  chest_pain      blood      serum_cholestoral  fasting_blo
    pressure
303 0.000000  303 0.000000  303 0.000000  303 0.000000  303 0.000000  303 0.000000
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   age                                  303 non-null    float64
 1   sex                                  303 non-null    float64
 2   chest_pain                          303 non-null    float64
 3   blood pressure                      303 non-null    float64
 4   serum_cholestoral                  303 non-null    float64
 5   fasting_blood_sugar                303 non-null    float64
 6   electrocardiographic               303 non-null    float64
 7   max_heart_rate                     303 non-null    float64
 8   induced_angina                     303 non-null    float64
 9   ST_depression                      303 non-null    float64
10  slope                              303 non-null    float64
11  vessels                             303 non-null    object
12  thal                                303 non-null    object
13  diagnosis                           303 non-null    int64
dtypes: float64(11), int64(1), object(2)
memory usage: 33.3+ KB

```

```

# checking for null
df.isna().sum()

```

```

age          0
sex          0
chest_pain   0
blood pressure 0
serum_cholestoral 0
fasting_blood_sugar 0
electrocardiographic 0
max_heart_rate 0
induced_angina 0
ST_depression 0
slope        0
vessels      0
thal         0
diagnosis    0
dtype: int64

```

```

## vessel and thal variables have some missing values shown as '?'
(df == '?').sum()

```

```

age          0
sex          0
chest_pain   0
blood pressure 0
serum_cholestoral 0
fasting_blood_sugar 0
electrocardiographic 0

```

```
max_heart_rate      0
induced_angina      0
ST_depression       0
slope               0
vessels             4
thal                2
diagnosis           0
dtype: int64
```

```
df['vessels'].value_counts()
```

```
0.0    176
1.0     65
2.0     38
3.0     20
?         4
Name: vessels, dtype: int64
```

```
df['thal'].value_counts()
```

```
3.0    166
7.0    117
6.0     18
?         2
Name: thal, dtype: int64
```

```
## removing the '?' missing values
```

```
df[df['vessels'] == '?']['vessels']
```

```
166    ?
192    ?
287    ?
302    ?
Name: vessels, dtype: object
```

```
df[df['thal'] == '?']['thal']
```

```
87     ?
266    ?
Name: thal, dtype: object
```

```
missing_vessels = df[df['vessels'] == '?']['vessels'].index
missing_thals = df[df['thal'] == '?']['thal'].index
```

```
df.drop(missing_vessels,inplace=True)
df.drop(missing_thals,inplace=True)
```

```
## After removing 6 missing values
df
```

	age	sex	chest_pain	blood pressure	serum_cholesterol	fasting_blood_sugar	electro
<b>0</b>	63.0	1.0	1.0	145.0	233.0	1.0	
<b>1</b>	67.0	1.0	4.0	160.0	286.0	0.0	
<b>2</b>	67.0	1.0	4.0	120.0	229.0	0.0	
<b>3</b>	37.0	1.0	3.0	130.0	250.0	0.0	
<b>4</b>	41.0	0.0	2.0	130.0	204.0	0.0	
...	...	...	...	...	...	...	
<b>297</b>	57.0	0.0	4.0	140.0	241.0	0.0	
<b>298</b>	45.0	1.0	1.0	110.0	264.0	0.0	
<b>299</b>	68.0	1.0	4.0	144.0	193.0	1.0	
<b>300</b>	57.0	1.0	4.0	130.0	131.0	0.0	
<b>301</b>	57.0	0.0	2.0	130.0	236.0	0.0	

297 rows × 14 columns

```
## Checking for age distribution
plt.figure(figsize=(12,8))
sns.distplot(df['age'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
```


```
`distplot` is a deprecated function and will be removed in a future version. Please
```

```
<matplotlib.axes. subplots.AxesSubplot at 0x7fdc6b89c590>
```

```
### age range of people with their gender 1: Male 0: Female
```

```
plt.figure(figsize=(10,6))
```

```
px.histogram(df,'age',color='sex')
```



```
<Figure size 720x432 with 0 Axes>
```

```
### Target distribution
```

```
df['diagnosis'].value_counts()
```

```
0    160
```

```
1     54
```

```
3     35
```

```
2     35
```

```
4     13
```

```
Name: diagnosis, dtype: int64
```

```
# ''' As given in the Dataset :The "goal" field refers to the presence of heart disease
#      in the patient. It is integer valued from 0 (no presence) to 4.
#      Experiments with the Cleveland database have concentrated on simply
#      attempting to distinguish presence (values 1,2,3,4) from absence (value
#      0). '''
```


```
df['Target'] = df['diagnosis'].apply(lambda x : 1 if x >= 1 else 0)
```

```
df
```

	age	sex	chest_pain	blood pressure	serum_cholesterol	fasting_blood_sugar	electrocardiogram
<b>0</b>	63.0	1.0	1.0	145.0	233.0	1.0	1.0
<b>1</b>	67.0	1.0	4.0	160.0	286.0	0.0	0.0
<b>2</b>	67.0	1.0	4.0	120.0	229.0	0.0	0.0
<b>3</b>	37.0	1.0	3.0	130.0	250.0	0.0	0.0
<b>4</b>	41.0	0.0	2.0	130.0	204.0	0.0	0.0
...	...	...	...	...	...	...	...
<b>297</b>	57.0	0.0	4.0	140.0	241.0	0.0	0.0
<b>298</b>	45.0	1.0	1.0	110.0	264.0	0.0	0.0
<b>299</b>	68.0	1.0	4.0	144.0	193.0	1.0	1.0
<b>300</b>	57.0	1.0	4.0	130.0	131.0	0.0	0.0
<b>301</b>	57.0	0.0	2.0	130.0	236.0	0.0	0.0

297 rows × 8 columns

```
#### Distribution of Diagnosis to sex
plt.figure(figsize=(10,6))
px.histogram(df,x='diagnosis',color='sex',barmode="group")
```




```
### Finding correlation of variables
df.corr()
```

	age	sex	chest_pain	blood pressure	serum_cholesterol	fasting_blood_sugar
<b>age</b>	1.000000	-0.092399	0.110471	0.290476	0.202644	
<b>sex</b>	-0.092399	1.000000	0.008908	-0.066340	-0.198089	
<b>chest_pain</b>	0.110471	0.008908	1.000000	-0.036980	0.072088	
<b>blood pressure</b>	0.290476	-0.066340	-0.036980	1.000000	0.131536	
<b>serum_cholesterol</b>	0.202644	-0.198089	0.072088	0.131536	1.000000	
<b>fasting_blood_sugar</b>	0.132062	0.038850	-0.057663	0.180860	0.012708	
<b>electrocardiographic</b>	0.149917	0.033897	0.063905	0.149242	0.165046	
<b>max_heart_rate</b>	-0.394563	-0.060496	-0.339308	-0.049108	-0.000075	
<b>induced_angina</b>	0.096489	0.143581	0.377525	0.066691	0.059339	
<b>ST_depression</b>	0.197123	0.106567	0.203244	0.191243	0.038596	
<b>slope</b>	0.159405	0.033345	0.151079	0.121172	-0.009215	
<b>diagnosis</b>	0.222156	0.226797	0.404248	0.159620	0.066448	
<b>Target</b>	0.227075	0.278467	0.408945	0.153490	0.080285	

```
plt.figure(figsize=(12,8))
sns.heatmap(df.corr())
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fdc66cb9950>



```
df['diagnosis']
```

```
0      0
1      2
2      1
3      0
4      0
..
297    1
298    1
299    2
300    3
301    1
Name: diagnosis, Length: 297, dtype: int64
```

```
### Train test split
```

```
X= df.drop(['diagnosis','Target'],axis=1)
```

```
y = df['Target']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
## Model Classification
```

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
### Taking average of 10 training examples
```



```

dt_avg = []
for i in range(10):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
    clf1 = DecisionTreeClassifier()
    dt_model = clf1.fit(X_train,y_train)
    dt_pred = dt_model.predict(X_test)
    dt_avg.append(accuracy_score(y_test,dt_pred))

print('Average Decision Tree Test accuracy:',sum(dt_avg)/10)

Average Decision Tree Test accuracy: 0.7333333333333334

from sklearn.metrics import classification_report, confusion_matrix,accuracy_score,log_loss
print(classification_report(y_test,dt_pred))

print(confusion_matrix(y_test,dt_pred))

```

	precision	recall	f1-score	support
0	0.77	0.75	0.76	36
1	0.64	0.67	0.65	24
accuracy			0.72	60
macro avg	0.71	0.71	0.71	60
weighted avg	0.72	0.72	0.72	60

```

[[27  9]
 [ 8 16]]

print('Log loss for Decision Tree model:',log_loss(y_test,dt_pred))

Log loss for Decision Tree model: 8.059114458598343

```